

In [1]:

```
import keras
import numpy as np
from keras.applications import vgg16
from keras.models import Sequential, Model
from keras.layers import Dense, Flatten, Dropout
from keras.preprocessing.image import ImageDataGenerator
from keras import optimizers
from keras.optimizers import Adam
import tensorflow as tf
gpus = tf.config.experimental.list_physical_devices('GPU')
if gpus:
    try:
        tf.config.experimental.set_virtual_device_configuration(
            gpus[0],
            [tf.config.experimental.VirtualDeviceConfiguration(memory_limit=2048)])
        logical_gpus = tf.config.experimental.list_logical_devices('GPU')
        print(len(gpus), "Physical GPUs,", len(logical_gpus), "Logical GPUs")
    except RuntimeError as e:

        print(e)
```

Using TensorFlow backend.

1 Physical GPUs, 1 Logical GPUs

In [2]:

```
train_data_dir = "D:/Dataset/Train"
validation_data_dir = "D:/Dataset/Test"
batch_size = 13
img_height = 208
img_width = 176
numClasses=4
epoch=13
train_datagen= ImageDataGenerator(
    rescale= 1./255,
    shear_range= 0.2,
    zoom_range= 0.2,
    horizontal_flip= True)

validation_datagen= ImageDataGenerator(
    rescale= 1./255)
train_generator= train_datagen.flow_from_directory(
    train_data_dir, target_size= (img_height, img_width),
    batch_size= batch_size,
    class_mode='categorical')
validation_generator= validation_datagen.flow_from_directory(
    validation_data_dir, target_size= (img_height, img_width),
    batch_size= batch_size,
    class_mode= 'categorical')

base_model= keras.applications.vgg16.VGG16(weights= 'imagenet', include_top= False, input_shape=(im
g_height, img_width, 3

))
```

Found 5472 images belonging to 4 classes.

Found 1279 images belonging to 4 classes.

In [3]:

```
add_model= Sequential()
add_model.add(Flatten(input_shape=base_model.output_shape[1:]))
add_model.add(Dense(250, activation= 'relu'))
add_model.add(Dropout(0.2))
add_model.add(Dense(4, activation='softmax'))
model= Model(inputs= base_model.input, outputs=add_model(base_model.output))
model.compile(loss='categorical_crossentropy', optimizer=optimizers.Adam(learning_rate=0.0001,
```

```

beta_1=0.9, beta_2=0.999, epsilon=1e-05, amsgrad=False),
    metrics=['accuracy'])
model.summary()
steps_per_epoch = train_generator.n // batch_size
validation_steps = validation_generator.n // batch_size

```

Model: "model_1"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 208, 176, 3)	0
block1_conv1 (Conv2D)	(None, 208, 176, 64)	1792
block1_conv2 (Conv2D)	(None, 208, 176, 64)	36928
block1_pool (MaxPooling2D)	(None, 104, 88, 64)	0
block2_conv1 (Conv2D)	(None, 104, 88, 128)	73856
block2_conv2 (Conv2D)	(None, 104, 88, 128)	147584
block2_pool (MaxPooling2D)	(None, 52, 44, 128)	0
block3_conv1 (Conv2D)	(None, 52, 44, 256)	295168
block3_conv2 (Conv2D)	(None, 52, 44, 256)	590080
block3_conv3 (Conv2D)	(None, 52, 44, 256)	590080
block3_pool (MaxPooling2D)	(None, 26, 22, 256)	0
block4_conv1 (Conv2D)	(None, 26, 22, 512)	1180160
block4_conv2 (Conv2D)	(None, 26, 22, 512)	2359808
block4_conv3 (Conv2D)	(None, 26, 22, 512)	2359808
block4_pool (MaxPooling2D)	(None, 13, 11, 512)	0
block5_conv1 (Conv2D)	(None, 13, 11, 512)	2359808
block5_conv2 (Conv2D)	(None, 13, 11, 512)	2359808
block5_conv3 (Conv2D)	(None, 13, 11, 512)	2359808
block5_pool (MaxPooling2D)	(None, 6, 5, 512)	0
sequential_1 (Sequential)	(None, 4)	3841254
Total params: 18,555,942		
Trainable params: 18,555,942		
Non-trainable params: 0		

In [4]:

```

model.fit_generator(train_generator,
    epochs=epoch,
    validation_data=validation_generator,
    verbose=1,
    steps_per_epoch=steps_per_epoch,
    validation_steps=validation_steps
)

```

Epoch 1/13

420/420 [=====] - 231s 551ms/step - loss: 1.1421 - accuracy: 0.4774 - val_loss: 0.9145 - val_accuracy: 0.5487

Epoch 2/13

420/420 [=====] - 210s 500ms/step - loss: 0.8866 - accuracy: 0.5877 - val_loss: 1.3873 - val_accuracy: 0.6003

Epoch 3/13

420/420 [=====] - 233s 556ms/step - loss: 0.7629 - accuracy: 0.6382 - val_loss: 0.6445 - val_accuracy: 0.6019

Epoch 4/13

```
420/420 [=====] - 264s 630ms/step - loss: 0.6504 - accuracy: 0.6912 - val
_loss: 1.1047 - val_accuracy: 0.6082
Epoch 5/13
420/420 [=====] - 271s 644ms/step - loss: 0.5715 - accuracy: 0.7485 - val
_loss: 0.5675 - val_accuracy: 0.6414
Epoch 6/13
420/420 [=====] - 269s 642ms/step - loss: 0.4776 - accuracy: 0.7941 - val
_loss: 1.8730 - val_accuracy: 0.6414
Epoch 7/13
420/420 [=====] - 245s 584ms/step - loss: 0.3844 - accuracy: 0.8403 - val
_loss: 1.4435 - val_accuracy: 0.6635
Epoch 8/13
420/420 [=====] - 219s 522ms/step - loss: 0.2834 - accuracy: 0.8877 - val
_loss: 3.6234 - val_accuracy: 0.6390
Epoch 9/13
420/420 [=====] - 246s 586ms/step - loss: 0.2033 - accuracy: 0.9221 - val
_loss: 1.0489 - val_accuracy: 0.6675
Epoch 10/13
420/420 [=====] - 260s 619ms/step - loss: 0.1671 - accuracy: 0.9397 - val
_loss: 2.8138 - val_accuracy: 0.6840
Epoch 11/13
420/420 [=====] - 245s 583ms/step - loss: 0.1510 - accuracy: 0.9419 - val
_loss: 1.3946 - val_accuracy: 0.6651
Epoch 12/13
420/420 [=====] - 231s 550ms/step - loss: 0.1194 - accuracy: 0.9564 - val
_loss: 2.6971 - val_accuracy: 0.6825
Epoch 13/13
420/420 [=====] - 232s 552ms/step - loss: 0.0962 - accuracy: 0.9639 - val
_loss: 0.9044 - val_accuracy: 0.7441
```

Out[4]:

```
<keras.callbacks.callbacks.History at 0x1f0970bb588>
```

In [5]:

```
model.save("model1.h5")
```

In []: